

Deployment options for OSGi applications in the cloud/edge

Deployment options for OSGi applications in the cloud/edge

Speaker



Dirk Fauth

*Research Engineer
Eclipse Committer*

ETAS GmbH
Borsigstraße 24
70469 Stuttgart

dirk.fauth@etas.com
www.etas.com


blog.vogella.com/author/fipro/
Twitter: [fipro78](#)

- 1. Deployment Variants**
- 2. Container**
- 3. Benchmark**
- 4. Conclusion**

Deployment Variants

Deployment Variants

General



Multiple JARs in a folder



Executable JAR



Custom JRE (jlink)

GraalVM™

Native Executable

Deployment Variants

Multiple JARs in a folder

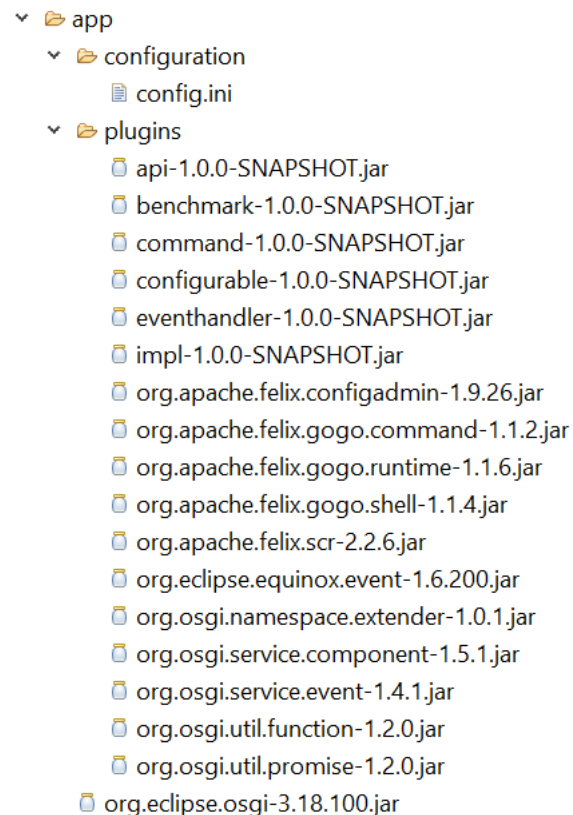
- Multiple JAR files (OSGi bundles) inside a folder
- Additional configuration file
- Launcher

```
org.eclipse.osgi
```

```
:org.eclipse.core.runtime.adaptor.EclipseStarter
```

```
java -jar org.eclipse.osgi-3.18.100.jar
```

- Build
 - maven-dependency-plugin
 - maven-resources-plugin



<https://www.eclipse.org/equinox/documents/quickstart-framework.php>

Deployment Variants

Executable JAR

- Executable JAR that includes each required bundle as embedded JAR file
- Configuration also included in the executable JAR
- Launcher

```
aQute.launcher.pre.EmbeddedLauncher
```

```
java -jar equinox-app.jar
```

- Build
 - bnd-maven-plugin
 - bnd-export-maven-plugin



<https://bnd.bndtools.org/>

<https://bndtools.org/>

<https://github.com/bndtools/bnd/tree/master/maven-plugins>

Deployment Variants

Custom JRE via jlink

- Create a custom JRE with `jlink` command of the JDK
 - *assemble and optimize a set of **modules** and their dependencies into a custom runtime image*

<https://docs.oracle.com/en/java/javase/17/docs/specs/man/jlink.html>

- Folder layout like JRE
- Launcher: `java` command

```
java [options] -m <module>[/<mainclass>]
```

```
/app/jre $ ls -l
total 20
drwxr-xr-x  2 appuser appuser 4096 Oct 14 08:37 bin
drwxr-xr-x  4 appuser appuser 4096 Oct 14 08:37 conf
drwxr-xr-x  9 appuser appuser 4096 Oct 14 08:37 legal
drwxr-xr-x  4 appuser appuser 4096 Oct 14 08:37 lib
-rw-r--r--  1 appuser appuser  140 Oct 14 08:37 release
/app/jre $
```

- Issue with OSGi and jlink
Most available OSGi bundles do not contain a `module-info.class`
→ **automatic module cannot be used with jlink**

JPMS


Native Executable with GraalVM

- *Native Image is a technology to compile Java code ahead-of-time to a binary – a native executable. A native executable includes only the code required at run time, that is the application classes, standard-library classes, the language runtime, and statically-linked native code from the JDK.*
- Can be created using the GraalVM `native-image` tool
 - From a **Class**, a **JAR (classpath)** or a **Module (modulepath)**
- “Closed world assumption”
 - all the bytecode in your application that can be called at run time must be known at build time
- Issue with OSGi and `native-image`
Dynamic classloading per bundle managed by OSGi Framework (Module Layer)
`java.lang.NullPointerException: A null service reference is not allowed.`

<https://www.graalvm.org/reference-manual/native-image/>

Deployment Variants

OSGi



Multiple JARs in a folder



Executable JAR



Custom JRE (jlink)



Native Executable

Deployment Variants

Custom JRE via jlink - OSGi

– Add `module-info.class`

– ModiTect

<https://github.com/moditect/moditect>

→ Intrusive change that adds an artifact to an existing published JAR

OSS license compatibility?

Checksum?

→ Requires knowledge on internals for generation

Maintenance?

– **Bndtools JPMS Support**

<https://bnd.bndtools.org/chapters/330-jpms.html>



Bndtools JPMS Support

Enable creation of `module-info.class` for each bundle, e.g. via `bnd-maven-plugin`

```
<plugin>
  <groupId>biz.aQute.bnd</groupId>
  <artifactId>bnd-maven-plugin</artifactId>
  <configuration>
    <bnd>
      <![CDATA[
Bundle-SymbolicName: ${project.groupId}.${project.artifactId}
-sources: true
-contract: *
-jpms-module-info:org.fipro.service.command;modules='org.apache.felix.configadmin'
-jpms-module-info-options: org.osgi.service.cm;ignore="true"
]]>
    </bnd>
  </configuration>
</plugin>
```

Enable creation of `module-info.class` for **executable jar** via `.bndrun` file

```
-jpms-module-info: \  
    ${project.groupId}.equinox.${project.artifactId};\  
    version=${project.version};\  
    ee=JavaSE-${java.specification.version}  
-jpms-module-info-options: jdk.unsupported;static=false
```

This makes the executable jar itself a module!

Deployment Variants

Custom JRE via jlink with Bndtools JPMS support

Build


```
$JAVA_HOME/bin/jlink \  
  --add-modules org.fipro.service.equinox.app \  
  --module-path equinox-app.jar \  
  --no-header-files \  
  --no-man-pages \  
  --output /app/jre
```

Launch

```
/app/jre/bin/java \  
  -m org.fipro.service.equinox.app/aQute.launcher.pre.EmbeddedLauncher
```

Deployment Variants

OSGi



+



Multiple JARs in a folder



+



Executable JAR



OSGi Connect

- *OSGi Connect allows for bundles to exist and be installed into the OSGi Framework from the flat class path, the module path (Java Platform Module System), a jlink image, or a native image.*

→ Allows to start an OSGi application without the full OSGi Module Layer

OSGi Core R8 – Connect Specification

<https://docs.osgi.org/specification/osgi.core/8.0.0/framework.connect.html>

Apache Felix Atomos

<https://github.com/apache/felix-atomos>

Ubiquitous OSGi - Android, Graal Substrate, Java Modules, Flat Class Path

<https://www.youtube.com/watch?v=KxmtzjHBumU>

OSGi R8, Felix 7, Atomos and the future of OSGi@Eclipse

<https://www.youtube.com/watch?v=oitFMbzt5s>

GraalVM Native Image with OSGi Connect

– Preparation

1. Add/use Atomos to be able to start the OSGi application from the flat classpath
2. Generate reachability metadata via tracing agent (reflection, resources, ...)
3. Update generated metadata

– Build



- Via GraalVM build plugins (Maven/Gradle)
- Docker multi-stage build using GraalVM container images

– Notes/Remarks


- `native-image` build only worked with **flat classpath** and **listing all jars explicitly**
- Build result is platform-dependent
- `atomos_lib` folder or index file needed for Atomos to discover bundles and load bundle entries
- Still not everything is working as expected (e.g. `scr:list` produces an empty output)

Deployment Variants

OSGi Connect / Apache Felix Atomos



Multiple JARs in a folder



Executable JAR



Custom JRE (jlink)

GraalVM™

Native Executable

Deployment (plain OSGi)

Multiple JARs in folder

Executable JAR

Custom JRE (jlink)

GraalVM Native Image

Deployment (OSGi Connect)

Multiple JARs in folder

Executable JAR

Custom JRE (jlink)

GraalVM Native Image

Container

“Size matters” – Find the right base image

Alpine vs. Debian vs. Ubuntu

Image	Size
alpine:3	5.54 MB
debian:bullseye-slim	80.50 MB
ubuntu:jammy	77.84 MB

Eclipse Temurin vs. IBM Semeru JDK vs. JRE

Image	Size
eclipse-temurin:17-jdk-jammy	~ 455 MB
eclipse-temurin:17-jdk-alpine	~ 356 MB
eclipse-temurin:17-jre-jammy	~ 266 MB
eclipse-temurin:17-jre-alpine	~ 168 MB
ibm-semeru-runtimes:open-17-jdk-jammy	~ 477 MB
ibm-semeru-runtimes:open-17-jre-jammy	~ 272 MB

Interlude: Distroless

- *"Distroless" images contain only your application and its runtime dependencies. They do not contain package managers, shells or any other programs you would expect to find in a standard Linux distribution.*

Image		Size
gcr.io/distroless/static-debian11	minimal Linux for "mostly-statically compiled" languages that do not require libc	2.36 MB
gcr.io/distroless/base-debian11	minimal Linux, glibc-based system	20.32 MB
gcr.io/distroless/java17-debian11	base image plus OpenJDK 17 and its dependencies	230.88 MB

- Distroless Java image is based on Debian and glibc, therefore bigger than an Alpine Temurin image
- Can be interesting in production for security reasons, but not for size

<https://github.com/GoogleContainerTools/distroless>

Java Best Practices

- Install only what you need
 - Use JRE instead of JDK
 - Use multi-stage builds (e.g. to create JRE or Native Image)
- Don't run Java apps as root
- Properly shutdown and handle events to terminate a Java application
- Take care of “container-awareness”

<https://snyk.io/blog/best-practices-to-build-java-containers-with-docker/>

<https://developers.redhat.com/articles/2022/04/19/java-17-whats-new-openjdk-container-awareness#>

<https://blog.openj9.org/2021/06/15/innovations-for-java-running-in-containers/>

Building Docker Images

- Use dedicated Docker files instead of generation tools
- Integrate image creation as part of the build via **fabric8io/docker-maven-plugin**
Maven/Gradle first

<https://github.com/fabric8io/docker-maven-plugin>
<http://dmp.fabric8.io/>

```
<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>docker-maven-plugin</artifactId>
  <extensions>true</extensions>
  <configuration>
    <images> ... </images>
  </configuration>
  <executions> ... </executions>
</plugin>
```

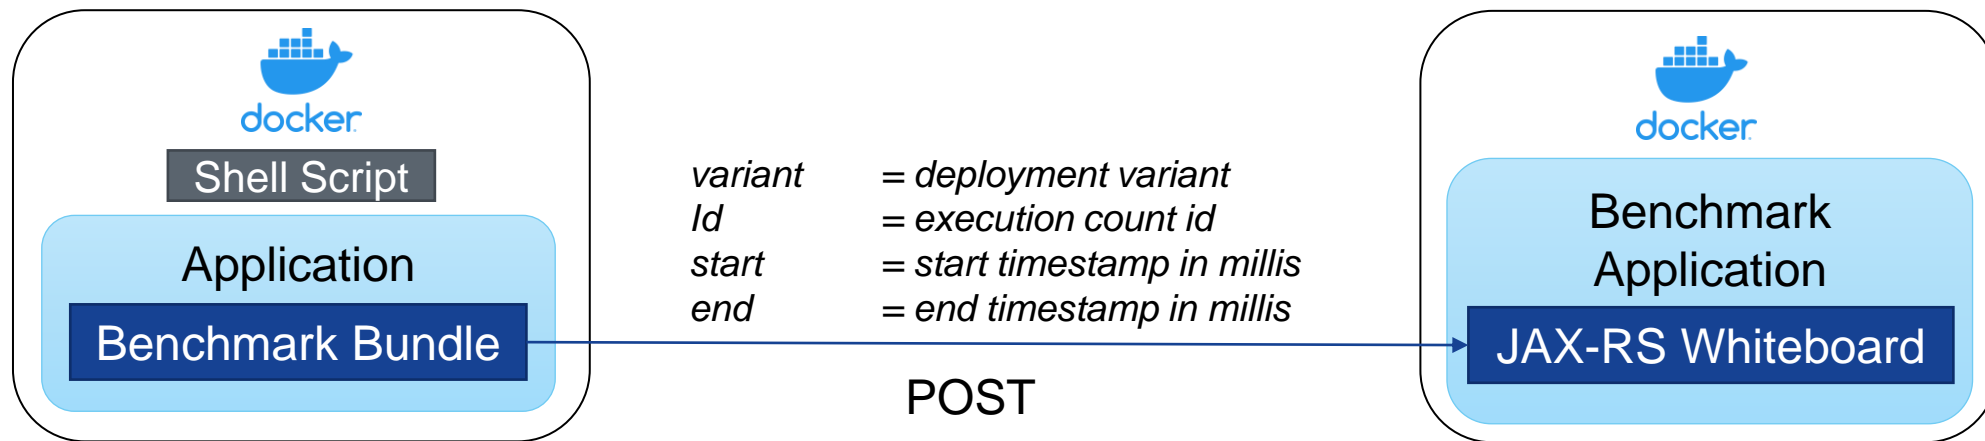
- Use multi-stage build to checkout sources and build in one container, then create new production container with build result only
Docker first

Deployment Variant – Base Image – Image Size

Deployment (plain OSGi)	Base Image	Size
Multiple JARs in folder	eclipse-temurin:17-jre-alpine	~ 171 MB
Executable JAR	eclipse-temurin:17-jre-alpine	~ 174 MB
Custom JRE (jlink)	alpine:3	~ 75 MB
Custom JRE (jlink/compressed)	alpine:3	~ 53 MB

Deployment (OSGi Connect)	Base Image	Size
Multiple JARs in folder	eclipse-temurin:17-jre-alpine	~ 171 MB
Custom JRE (jlink)	alpine:3	~ 75 MB
Custom JRE (jlink/compressed)	alpine:3	~ 53 MB
GraalVM Native Image	scratch	~ 38 MB
	alpine:3	~ 43 MB

Benchmark



Benchmark Bundle / Immediate Component

- Get start timestamp from system property
- Get current timestamp
- Send POST request via `java.net.http.HttpClient`
- Shutdown

Shell script

- Execute application multiple times in for-loop (clean/cache)
- Pass start timestamp as system property

Benchmark Images

Deployment Variant – Base Image – Image Size – Benchmark Image Size

Deployment (plain OSGi)	Base Image	Size	Size Benchmark
Multiple JARs in folder	eclipse-temurin:17-jre-alpine	~ 171 MB	~ 173 MB
Executable JAR	eclipse-temurin:17-jre-alpine	~ 174 MB	~ 176 MB
Custom JRE (jlink)	alpine:3	~ 75 MB	~ 78 MB
Custom JRE (jlink/compressed)	alpine:3	~ 53 MB	~ 55 MB

- + coreutils
- + nanosecond support
- + benchmark bundle
- + java.net.http module
- + shell script support

Deployment (OSGi Connect)	Base Image	Size	Size Benchmark
Multiple JARs in folder	eclipse-temurin:17-jre-alpine	~ 171 MB	~ 173 MB
Custom JRE (jlink)	alpine:3	~ 75 MB	~ 78 MB
Custom JRE (jlink/compressed)	alpine:3	~ 53 MB	~ 55 MB
GraalVM Native Image	scratch alpine:3	~ 38 MB (~ 43 MB)	(~ 46 MB) ~ 53 MB

Benchmark Results

Deployment (plain OSGi)	Base Image	Size	Size Benchmark	Startup Clean	Startup Cache
Multiple JARs in folder	eclipse-temurin:17-jre-alpine	~ 171 MB	~ 173 MB	~ 982 ms	~ 901 ms
Executable JAR	eclipse-temurin:17-jre-alpine	~ 174 MB	~ 176 MB	~ 1087 ms	~ 1099 ms
Custom JRE (jlink)	alpine:3	~ 75 MB	~ 78 MB	~ 1336 ms	~ 1345 ms
Custom JRE (jlink/compressed)	alpine:3	~ 53 MB	~ 55 MB	~ 1497 ms	~ 1505 ms

Deployment (OSGi Connect)	Base Image	Size	Size Benchmark	Startup Clean	Startup Cache
Multiple JARs in folder classpath modulepath	eclipse-temurin:17-jre-alpine	~ 171 MB	~ 173 MB	~ 1122 ms	~ 973 ms
				~ 1194 ms	~ 1052 ms
Custom JRE (jlink)	alpine:3	~ 75 MB	~ 78 MB	~ 1439 ms	~ 1326 ms
Custom JRE (jlink/compressed)	alpine:3	~ 53 MB	~ 55 MB	~ 1593 ms	~ 1445 ms
GraalVM Native Image	scratch	~ 38 MB	(~ 46 MB)	-	-
	alpine:3	(~ 43 MB)	~ 53 MB	~ 34 ms	-

Conclusion

- All Java deployment variants possible for OSGi applications via
 - Bndtools JPMS support
 - OSGi Connect (Felix Atomos)
- Different deployment variants have different startup & runtime behaviors
Consider JIT vs. AOT compilation
- Make decision about variant dependent on the use case,
e.g. short running executables in container vs. long running application servers
- Further optimizations possible by configuring the Java runtime,
e.g. Container-awareness, Garbage Collection, *Checkpoint & Restore*, etc.

“AOT like startup performance with JIT runtime performance and behaviour”

- **CRaC (Coordinated Restore at Checkpoint)**
 - OpenJDK CRaC JDK
 - Azul Zulu JDK with CRaC support
 - Enable via `-XX:CRaCCheckpointTo=<path>`
 - Create a checkpoint via `jdk.crac.Core` API
 - Create a checkpoint via `jcmd <PID> JDK.checkpoint` from separate shell process
- **Open Liberty InstantOn / OpenJ9 CRIU support**
 - IBM Semeru (OpenJDK/OpenJ9)
 - Enable via `-XX:+EnableCRIUSupport`
 - Creating a checkpoint via `org.eclipse.openj9.criu.CRIUSupport` API

Current issues

- **CRaC (Coordinated Restore at Checkpoint)**
 - Only Linux, currently no Alpine (musl) variant available
 - Building a custom JRE (jlink) currently not easily possible (need to manually copy criu to the system)
 - Containers with CRaC support not yet available out-of-the-box
- **Open Liberty InstantOn / OpenJ9 CRIU support**
 - Only Linux, no Alpine
 - Building a custom JRE (jlink)?
 - Creating a checkpoint only possible via API

– April 2023

Azul Zulu builds with CRaC support generally available for Java 17 on Linux

– June 2023

Open Liberty InstantOn out of beta with release 23.0.0.6

– June 2023

Oracle announced **Oracle GraalVM** distributed under **GraalVM Free License** with release 23.0

→ Oracle GraalVM Enterprise now available for free

→ Container images not yet available

https://github.com/fipro78/osgi_deployment_options

Thank you

Dirk Fauth

ETAS/ENA

dirk.fauth@etas.com